Shaping Knowledge Graphs
ISWC'24 Tutorial

Jose Emilio Labra Gayo
WESO Research group
University of Oviedo, Spain

About me…

Main researcher at WESO (Web Semantics Oviedo)

Some books:

- "*Web semántica*" (in Spanish), 2012
- "*Validating RDF data*", 2017
- "Knowledge Graphs", 2021

…and software:

- SHaclEX (Scala library, implements ShEx & SHACL)
- RDFShape (RDF playground)
- rudof (RDF & Shapes library in Rust)

http://labra.weso.es

2012

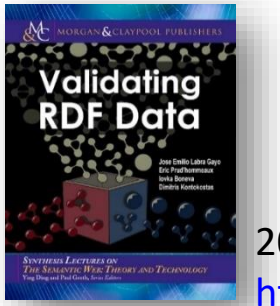

2017 HTML version:
http://book.validatingrdf.com

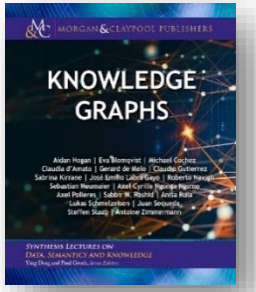

2021, HTML version
https://kgbook.org/

Contents

- Introduction to Knowledge graphs
- Types of Knowledge Graphs:
  - RDF, Property graphs, Wikibase, RDF-Star
- Validating RDF: ShEx & SHACL
- Validating Property Graphs
- Validating Wikibase and Wikidata graphs
- Validating RDF-Star
- Applications:
  - Inferring shapes from data, Knowledge Graphs Subsets, etc.

Knowledge Graphs

Current notion of Knowledge Graphs Popular after Google, 2012*

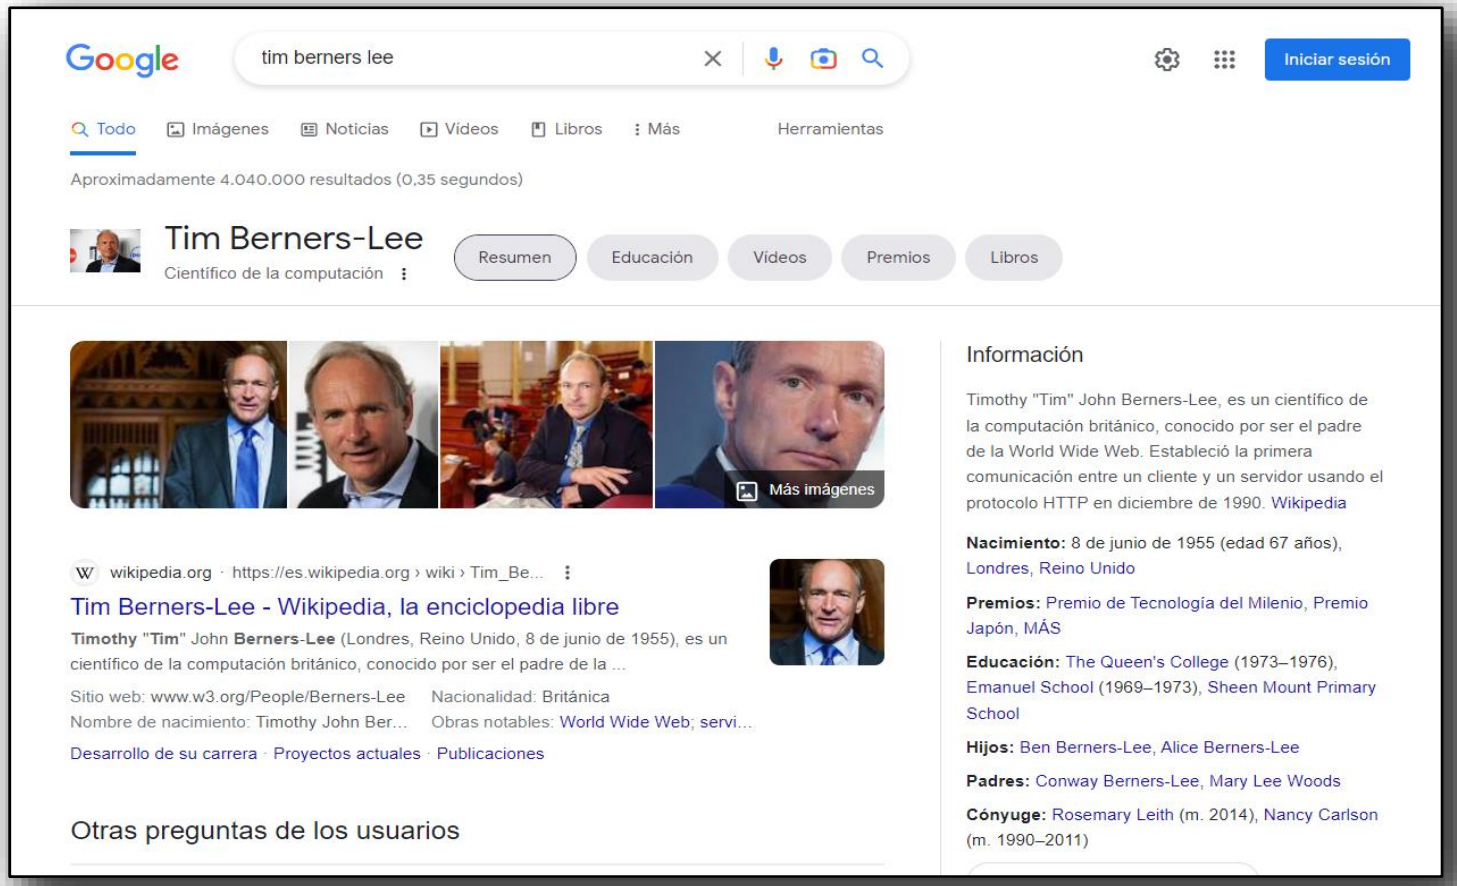

Link: https://www.blog.google/products/search/introducing-knowledge-graph-things-not/

Knowledge graphs

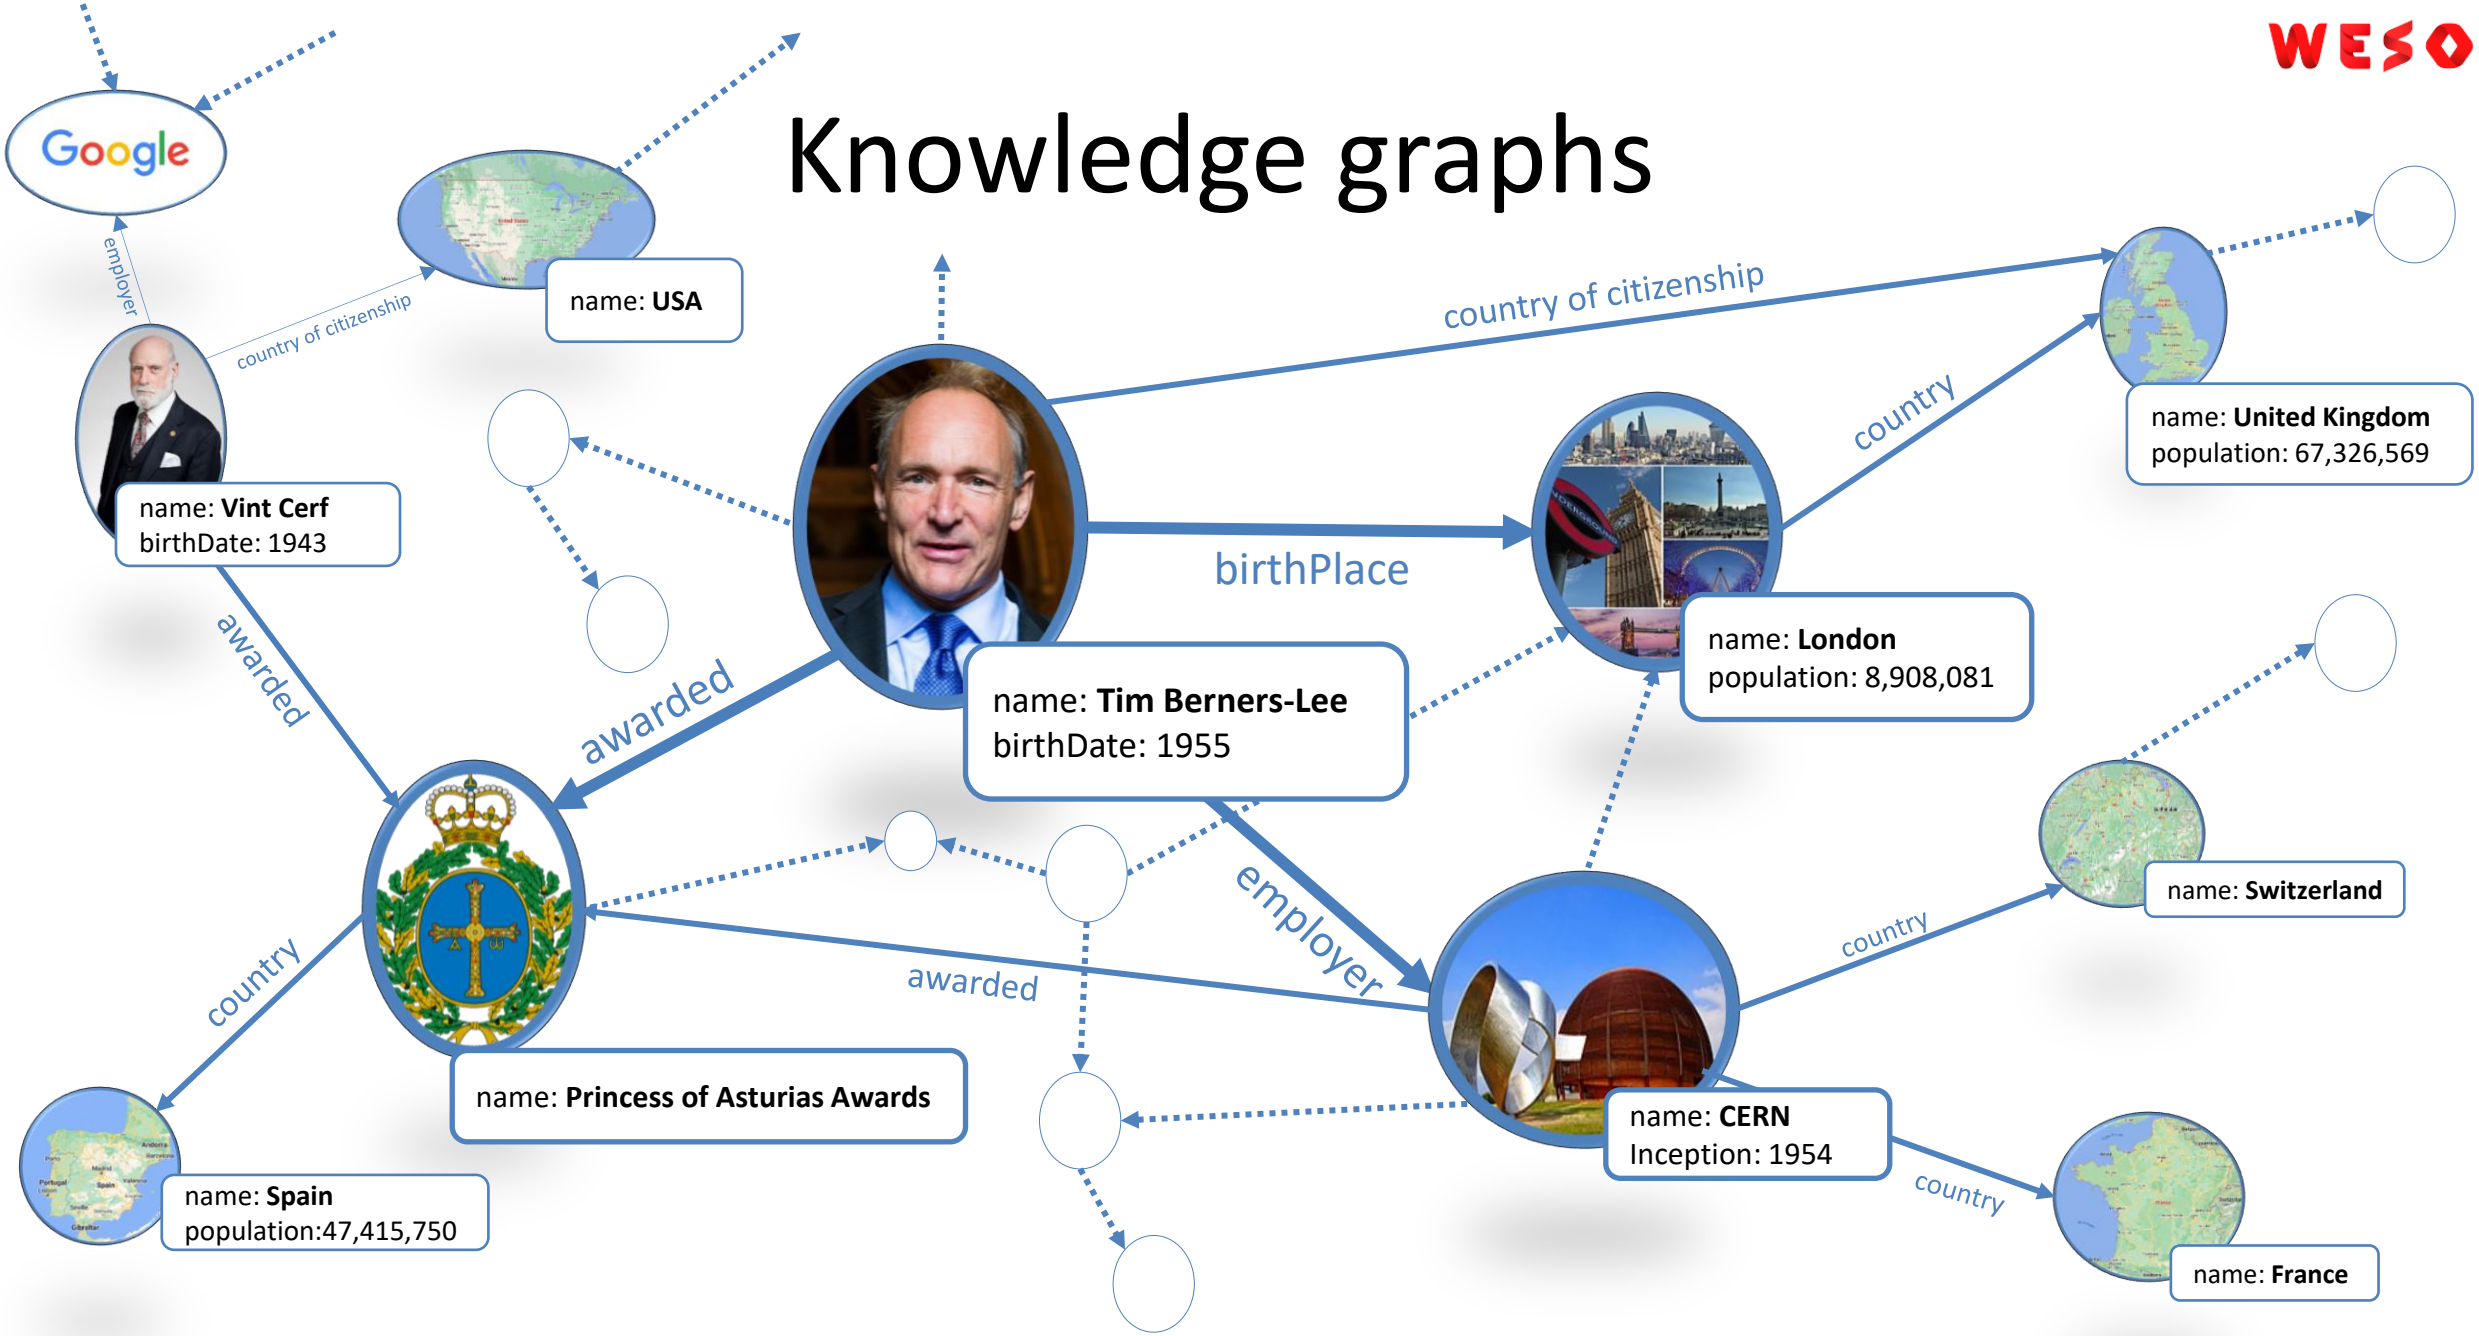

Knowledge graphs

Knowledge graph = *a **graph of data***

*intended to accumulate and convey **knowledge** of the real world*

*whose nodes represent **entities** of interest and*

*whose edges represent **relations** between these entities.*

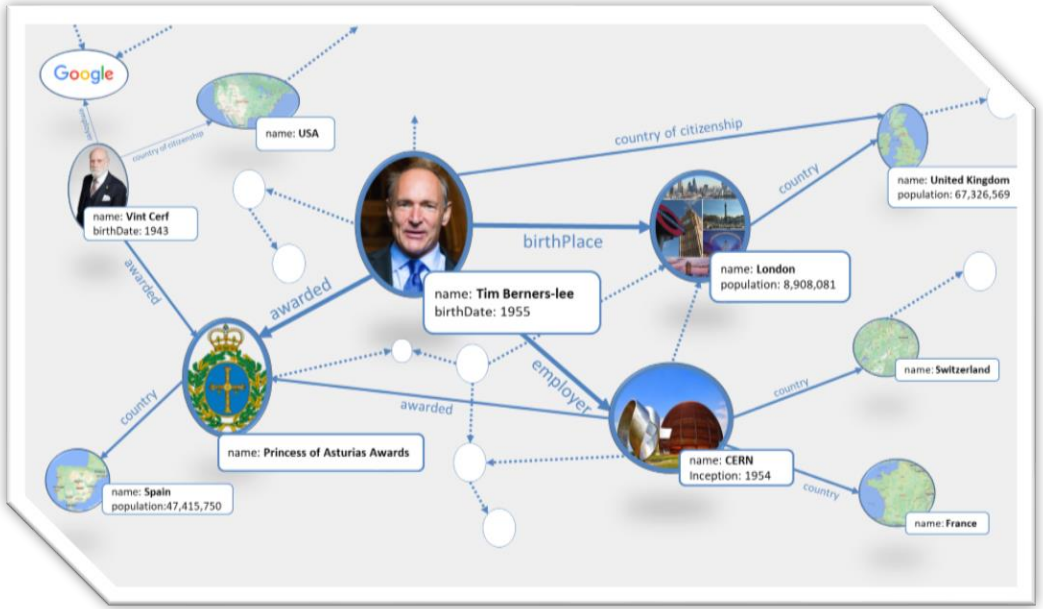

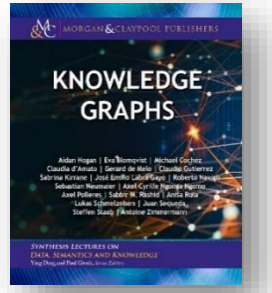

https://kgbook.org/

Applications of Knowledge graphs

Improve search results

Question answering

Data governance

Handling heterogenous data

Recommender systems

Chatbots and NLP

. . .

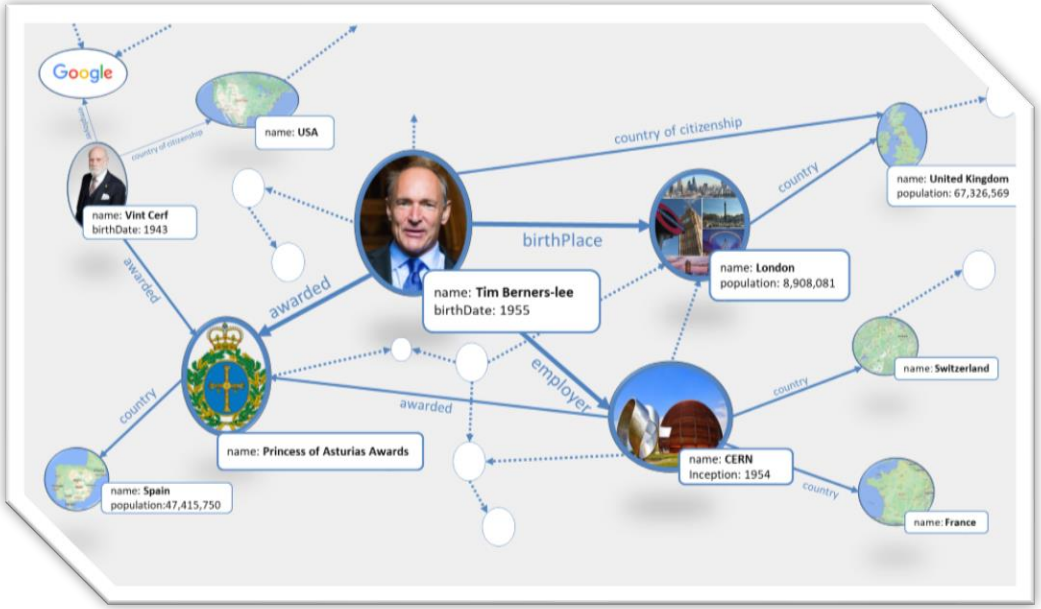

Contents

Introduction to Knowledge graphs

Types of Knowledge Graphs:

- RDF, Property graphs, Wikibase, RDF-Star

Validating RDF: ShEx & SHACL

Validating Property Graphs

Validating Wikibase and Wikidata graphs

Validating RDF-Star

Applications:

- Inferring shapes from data, Knowledge Graphs Subsets, etc.

Types of Knowledge graphs

- Open Knowledge graphs
  - Cross-domain: Wikidata, Dbpedia, Freebase, YAGO, ...
  - Domain specific
    - Academic: Open citations, SciGraph, Microsoft Academic Knowledge Graph, ...
    - Life sciences: UniProt, PubChem, PDB, ...
    - Government: EU Knowledge graph, ...
    - ...
- Enterprise Knowledge graphs
  - Web search: Google, Bing...
  - Commerce: AirBnb, Amazon, eBay, Uber,...
  - Social networks: Linkedin, Facebook,...
  - Finance: Banca d'Italia, Bloomberg, Wells Fargo, Capital One,...
  - ...

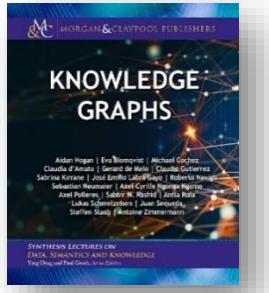

https://kgbook.org/

Wikidata as an example

Wikidata created in 2012 as a collaborative knowledge graph

- Initial goal:
  - Support multilingual infoboxes in Wikipedia

Wikidata as an example

English Wikipedia page of Tim Berners-Lee

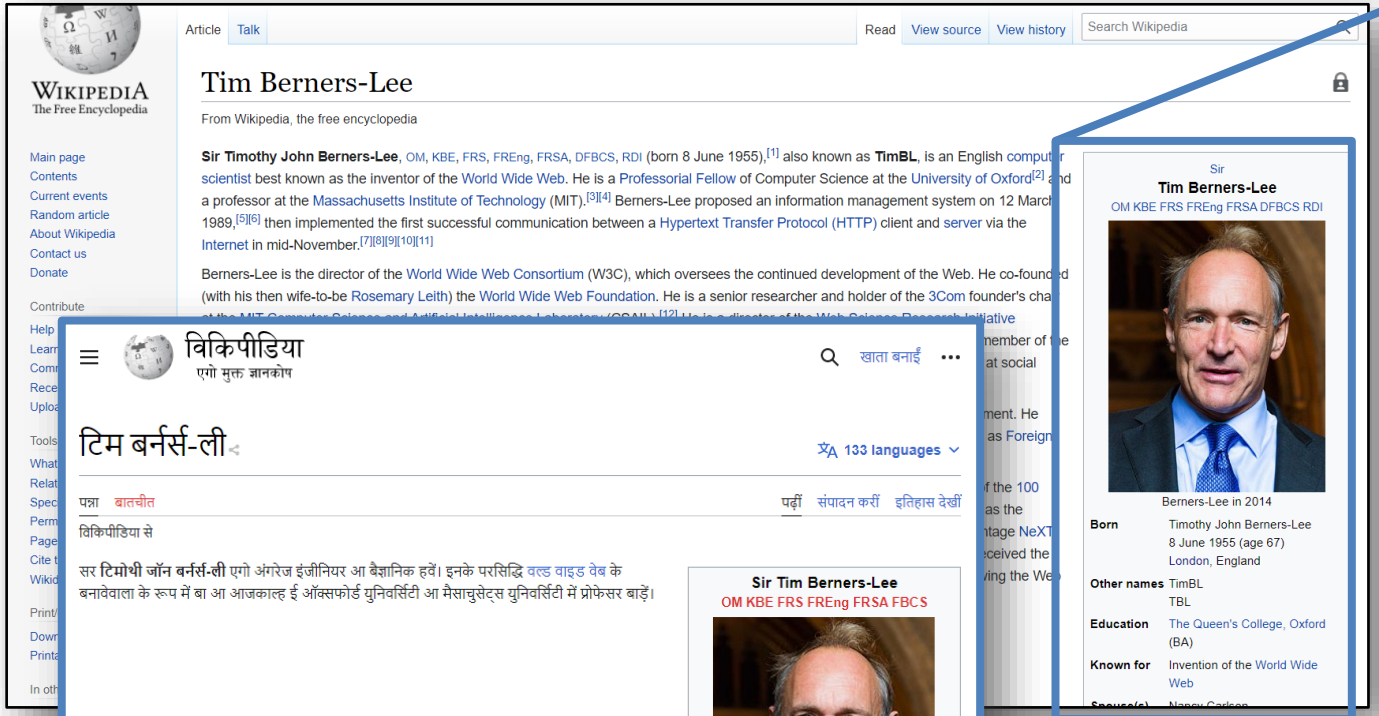

Bihari Wikipedia of Tim Berners-Lee

http://www.wikidata.org/entity/Q80

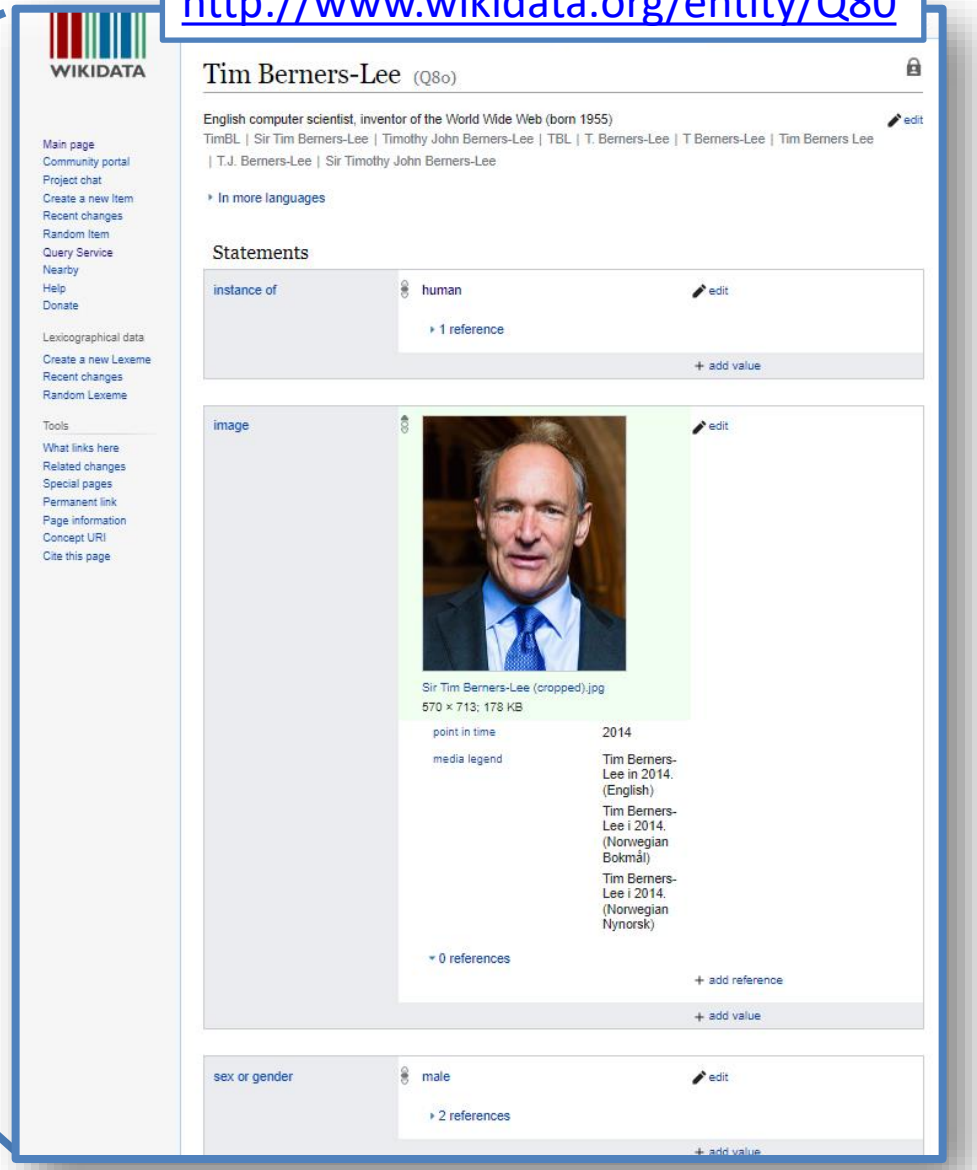

Wikidata: some features

Collaborative: anyone can edit

Free and open license

Currently (01/2023): 101m items, 1,8b edits

Co-edited by humans and bots: 23k active users, 343 bots

Open Wikidata Query Service: Public SPARQL endpoint

Dumps freely available: 109Gb compressed

Software that supports Wikidata = Wikibase

Evolution

Timeline with some concepts and technologies...

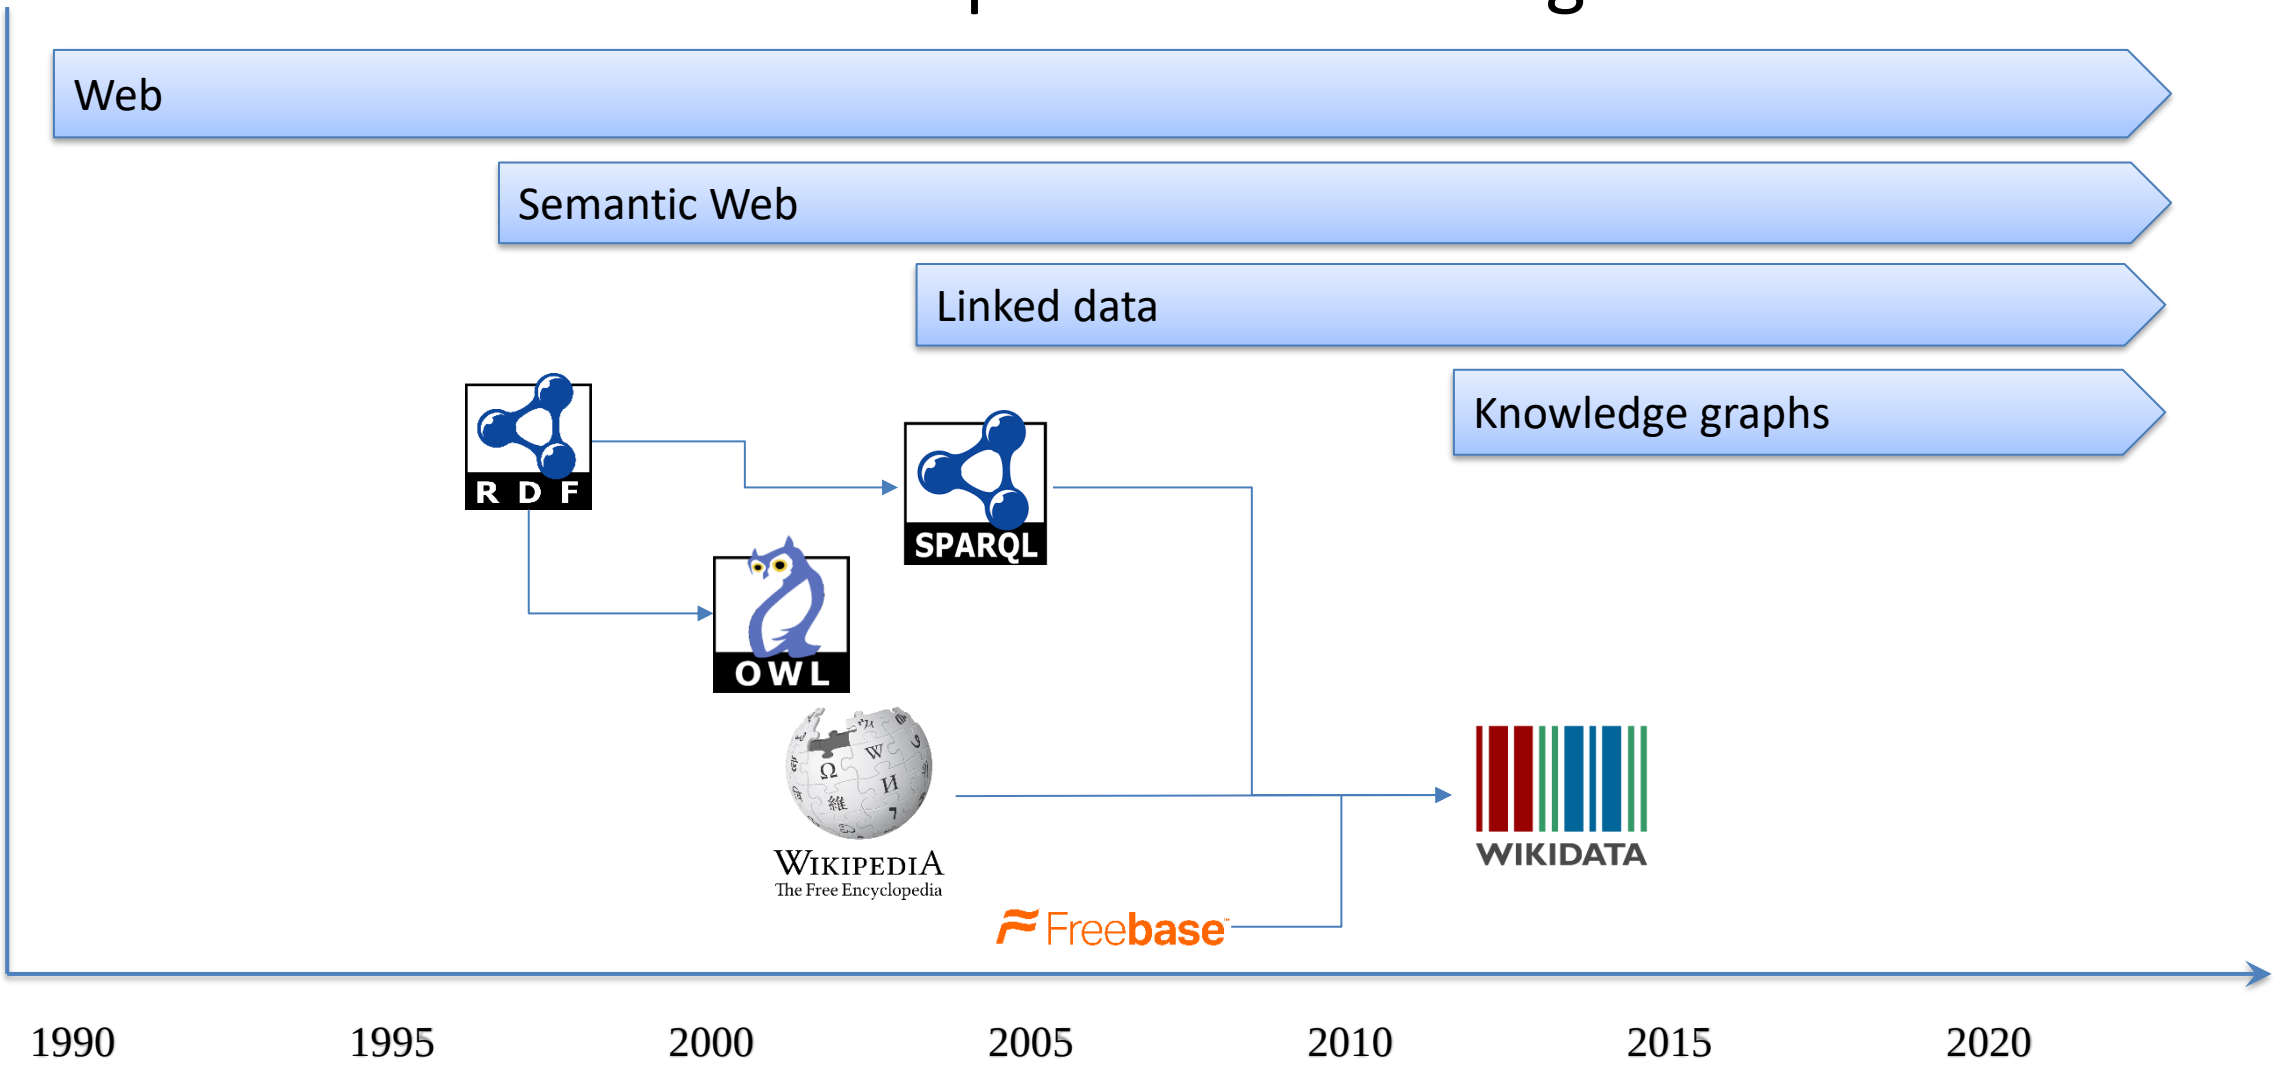

Knowledge Graphs models

3 popular knowledge graphs models

- RDF based
- Property graphs
- Wikibase graphs

RDF graphs

RDF = W3C recommendation (since 98)

Lingua franca of Semantic Web

- Based on triples
  - (subject, predicaje, object)
- Most nodes are URIs
  - Interoperability

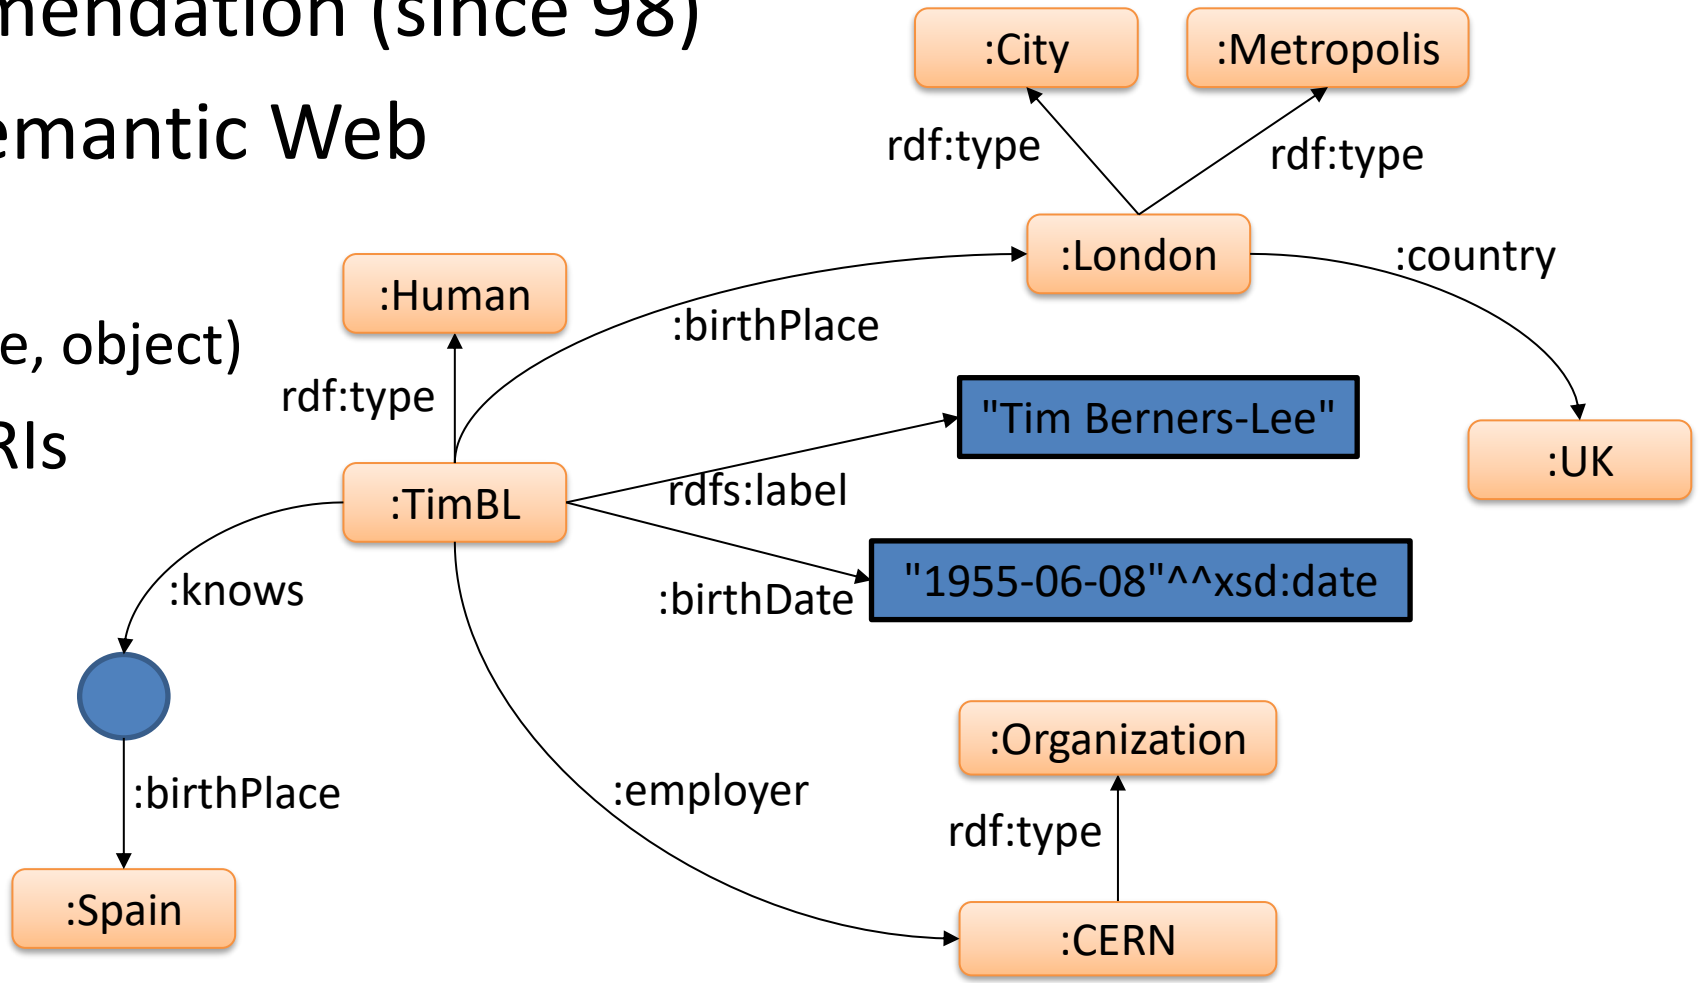

RDF ecosystem

One data model, several syntaxes: Turtle, N-Triples, JSON-LD

Vocabularies: RDF Schema, OWL, SKOS, etc.

Turtle

```
prefix :       <http://example.org/>
prefix rdfs:   <http://www.w3.org/2000/01/rdf-schema#>
prefix xsd:    <http://www.w3.org/2001/XMLSchema#>
prefix rdf:    <http://www.w3.org/1999/02/22-rdf-syntax-ns#>

:timbl  rdf:type     :Human ;
        :birthPlace  :london ;
        rdfs:label   "Tim Berners-Lee" ;
        :birthDate   "1955-06-08"^^xsd:date ;
        :employer    :CERN ;
        :knows       _:1    .
:london rdf:type     :City, :Metropolis ;
        :country     :UK .
:CERN   rdf:type     :Organization .
_:1     :birthPlace  :Spain .
```

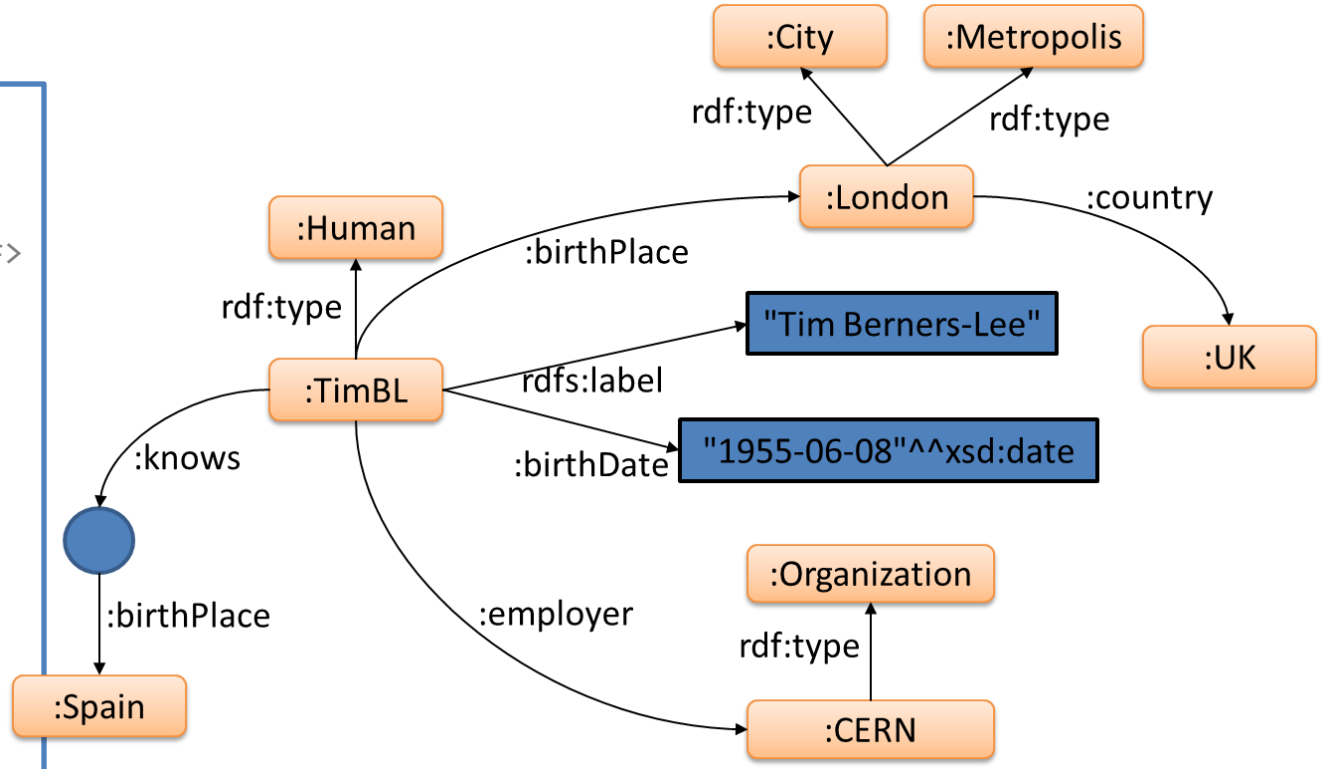

RDF ecosystem: SPARQL

SPARQL is an RDF query language and protocol

It enables the creation of SPARQL endpoints

```
select ?person ?date ?country where {
    ?person :birthDate  ?date .
    ?person :birthPlace ?p .
    ?p      :country    ?country
}
```

?person	?date	?country
:timbl	1955-06-08	:UK

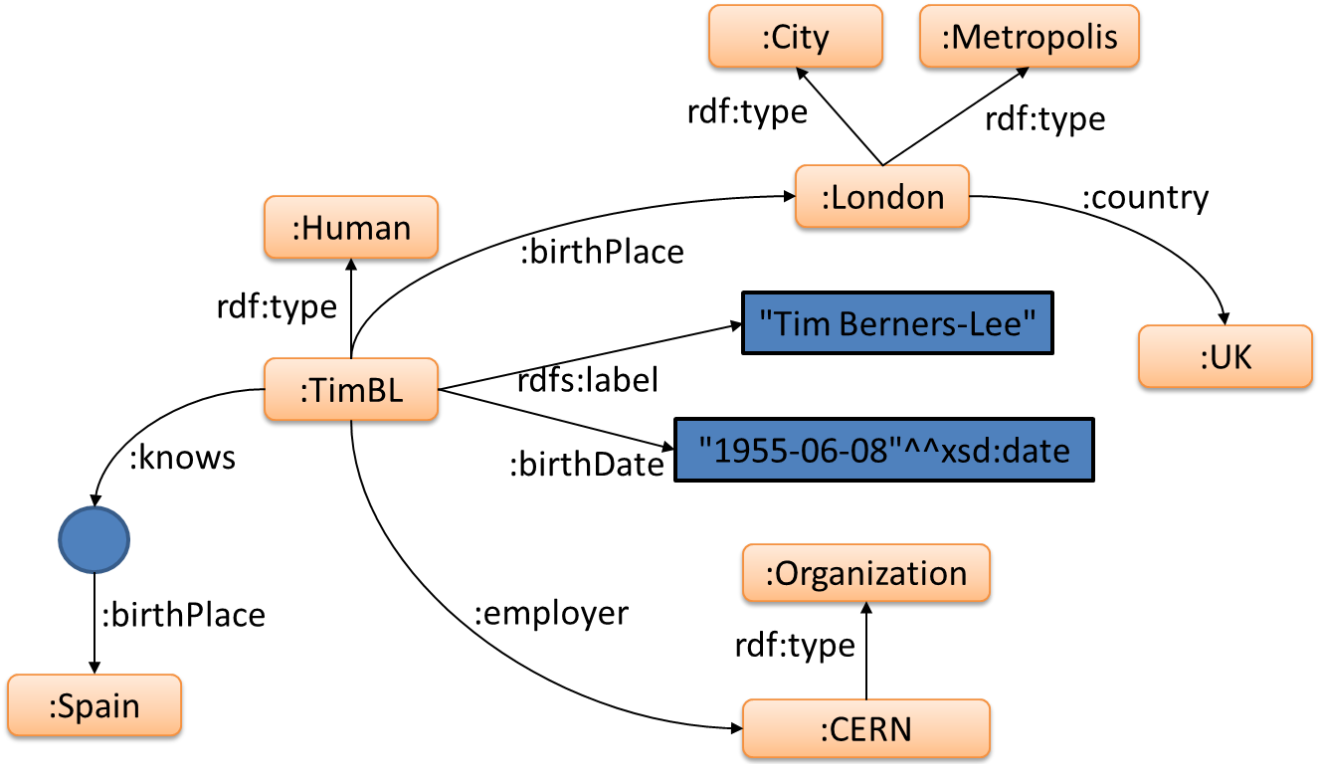

RDF1.2 (RDF-Star)

Triple terms

Add statements about triples

Reifiers

```
prefix :      <http://example.org/>

_:r1 rdf:reifies << :timbl :employer :CERN >> .
_:r1 :start "1980-06";
     :end   "1980-12" .
```

```
_:r1 rdf:reifies << :timbl :employer :CERN >> .
_:r2 :start "1984;
     :end   "1994" .
```

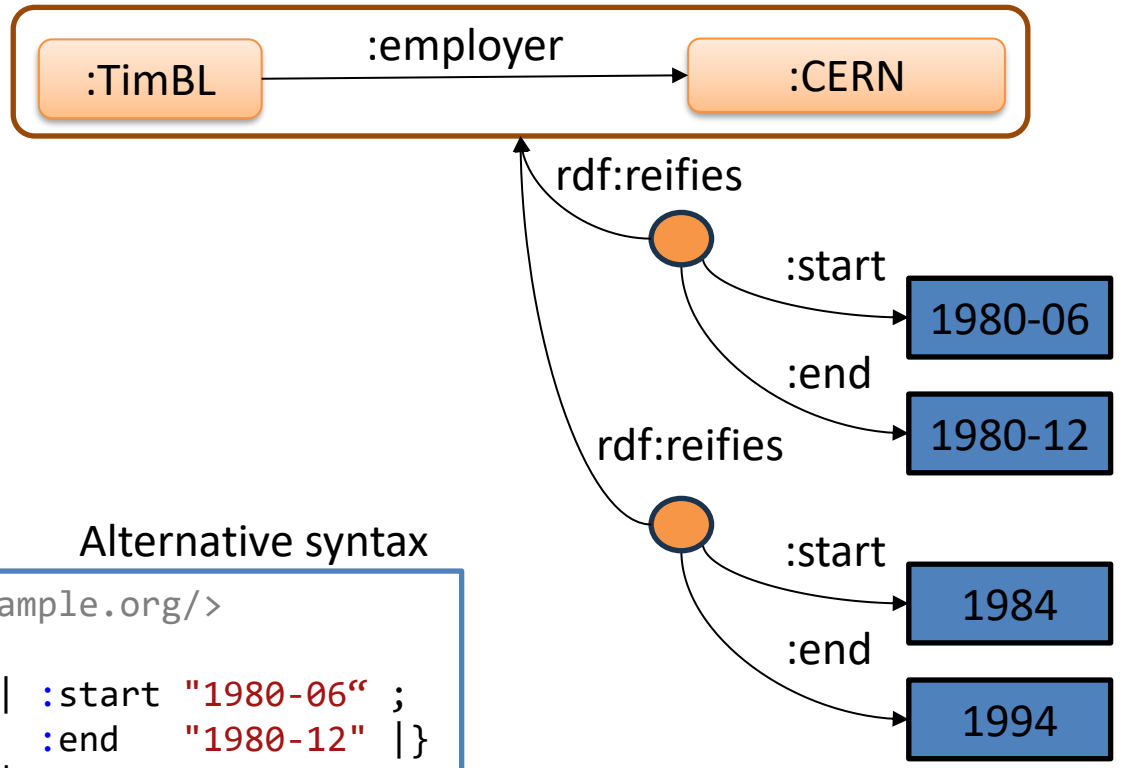

Alternative syntax

```
prefix :      <http://example.org/>

:timbl :employer :CERN {| :start "1980-06“ ;
                          :end   "1980-12" |}
                       {| :start "1984“ ;
                          :end   "1994" |} .
```

Property graphs

Popularized by graph databases

Example: Neo4j

Nodes and relations can be qualified

Nodes can have types

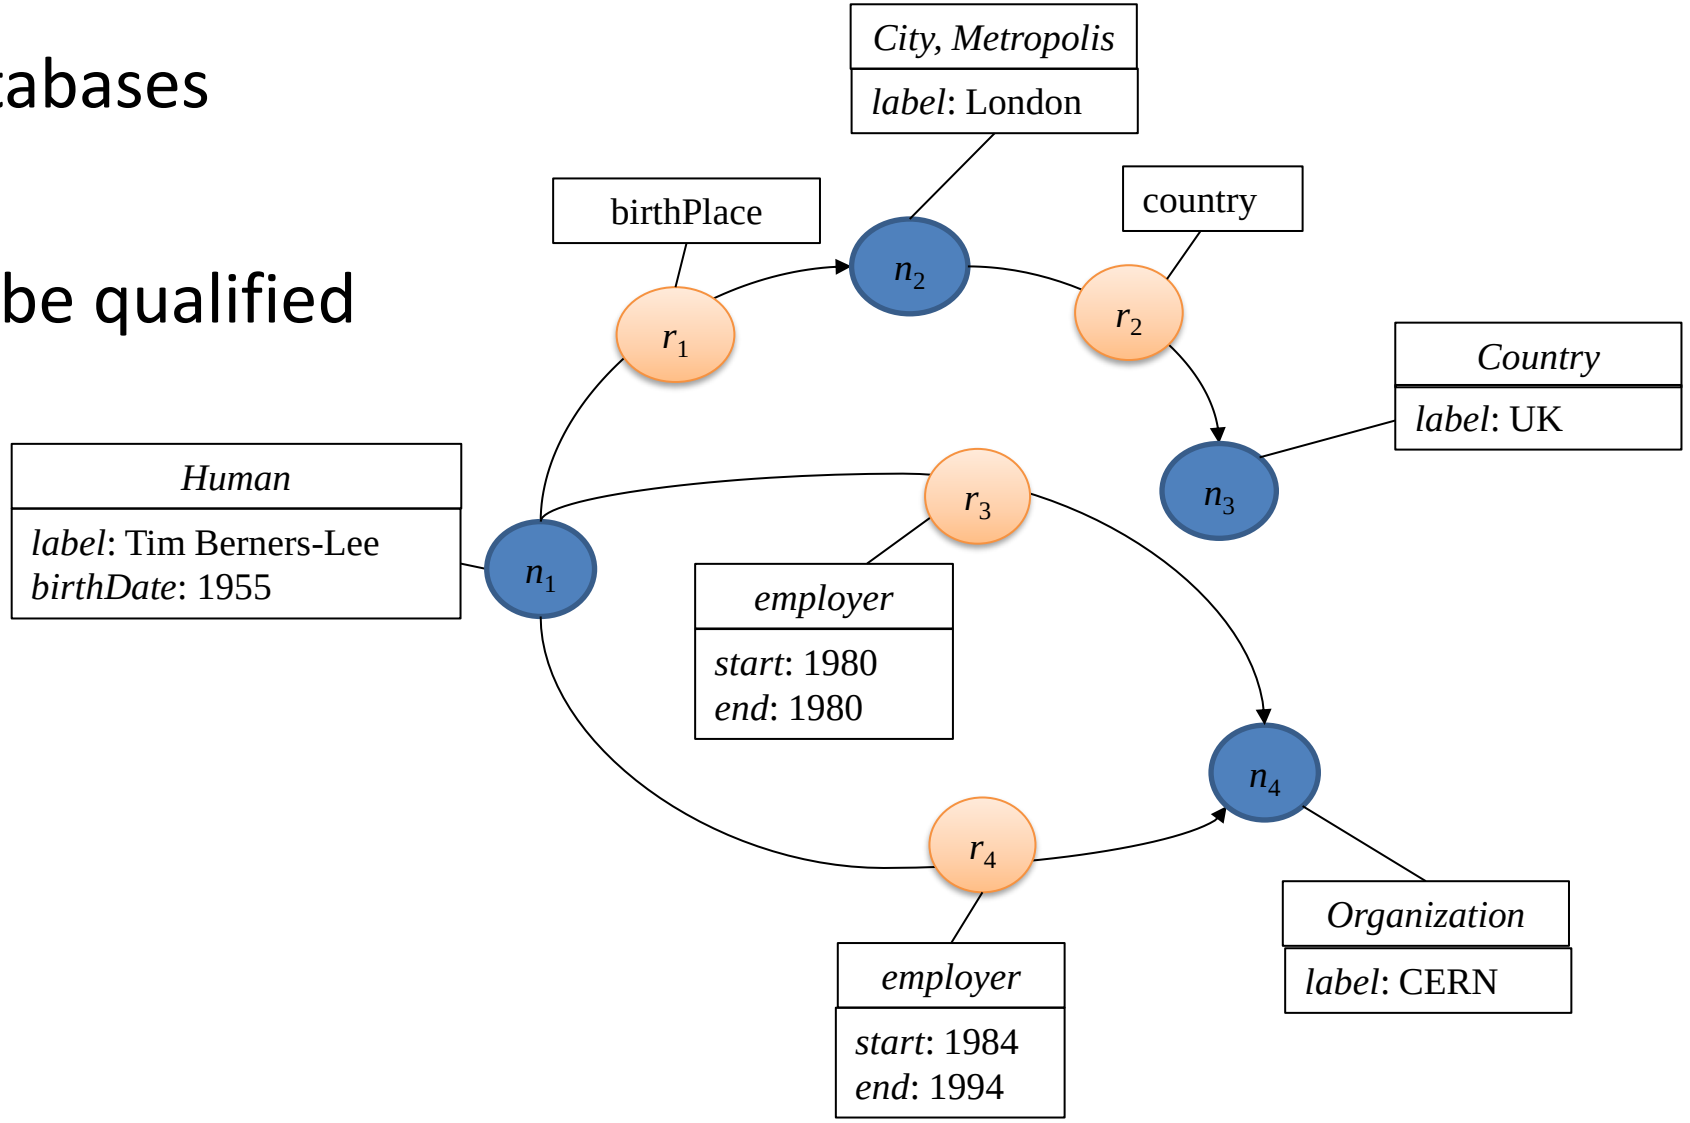

Wikibase graphs

Popularized by Wikidata
Wikibase = software supporting Wikidata
The values can be nodes in the graph
Example:
Tim Berners Lee
http://www.wikidata.org/entity/Q80

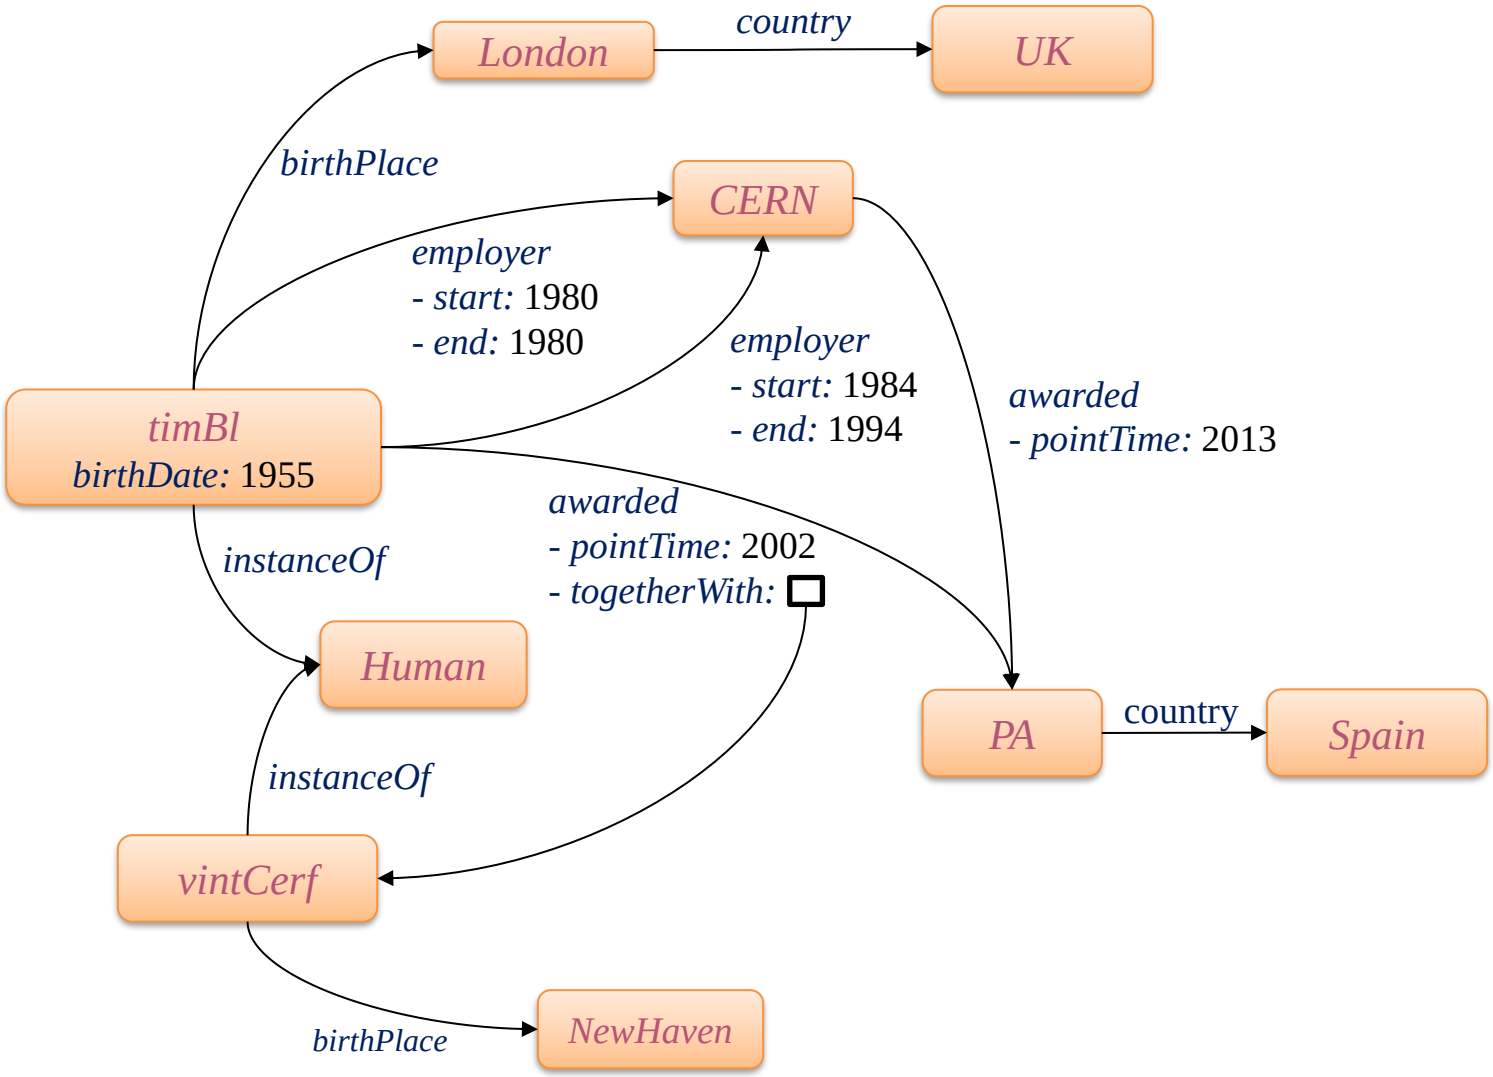

Wikibase graphs and SPARQL

Wikibase graphs generate RDF serializations for each item

SPARQL endpoint and Query service available

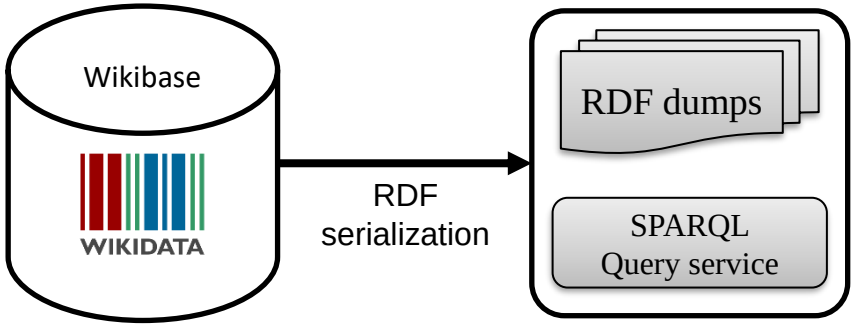


```
select ?name ?date?country where {
 wd:Q80 wdt:P1559 ?name .
 wd:Q80 wdt:P569  ?date .
 wd:Q80 wdt:P19   ?place .
 ?place wdt:P17   ?country
}
```

?name	?date	?country
Tim Berners-lee	1955-06-08	:UK

Try it: https://w.wiki/5yGu

Contents

Introduction to Knowledge graphs

Types of Knowledge Graphs:

RDF, Property graphs, Wikibase, RDF-Star

Validating RDF: ShEx & SHACL

Validating Property Graphs

Validating Wikibase and Wikidata graphs

Validating RDF-Star

Applications:

Inferring shapes from data, Knowledge Graphs Subsets, etc.

RDF Data Model

Overview of RDF Data Model and simple exercise

Link to slides about
RDF Data Model

https://doi.org/10.6084/m9.figshare.13174562

RDF, the good parts...

- RDF as an integration language
- RDF as a *lingua franca* for semantic web and linked data
- RDF flexibility
  - Data can be adapted to multiple environments
  - Open and reusable data by default
- RDF for knowledge representation
- RDF data stores & SPARQL

RDF, the other parts

Consuming & producing RDF

- Multiple serializations: Turtle, RDF/XML, JSON-LD, ...
- Embedding RDF in HTML
- Describing and validating RDF content

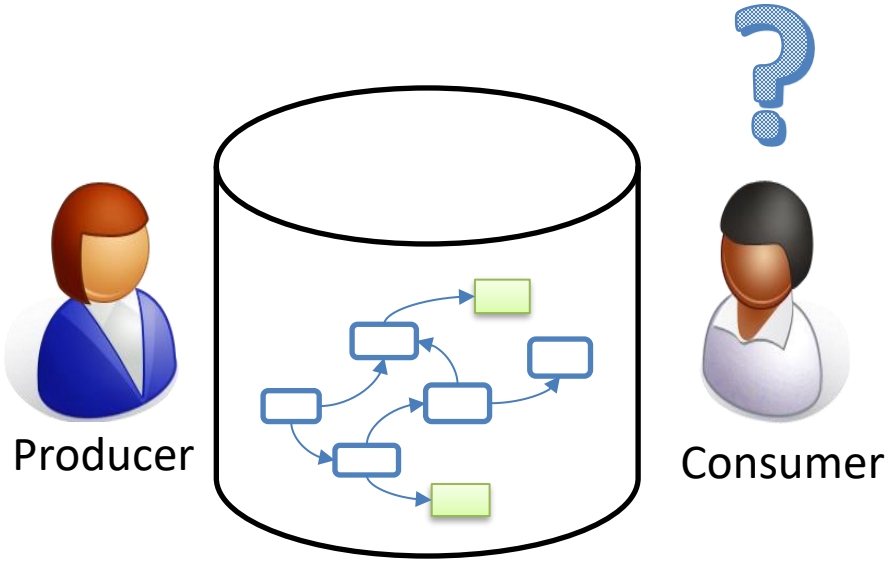

Why describe & validate RDF?

For producers

- Developers can understand the contents they are going to produce
- They can ensure they produce the expected structure
- Advertise and document the structure
- Generate interfaces

For consumers

- Understand the contents
- Verify the structure before processing it
- Query generation & optimization

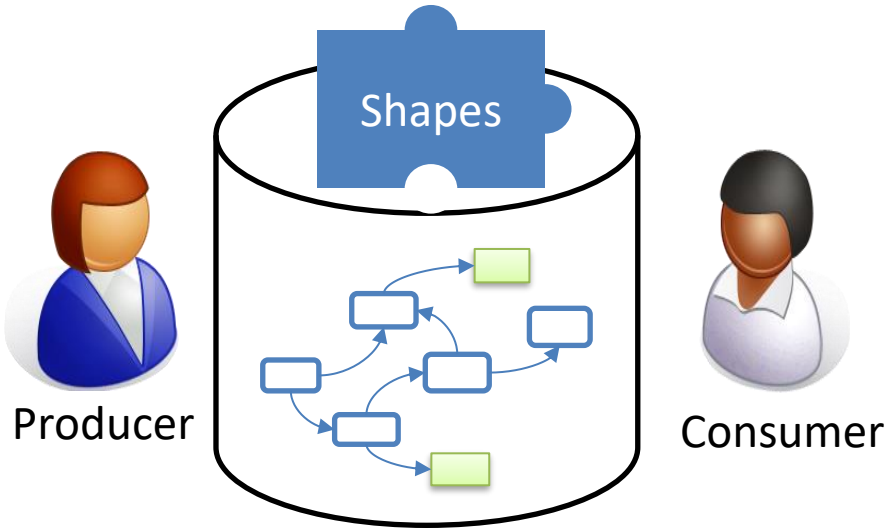

Similar technologies

Technology	Schema
Relational Databases	DDL
XML	DTD, XML Schema, RelaxNG, Schematron
Json	Json Schema
RDF	?

Fill that gap

Understanding the problem

RDF is composed by nodes and arcs between nodes

We can describe/check

- The form of the node itself (node constraint)
- The number of possible arcs incoming/outgoing from a node
- The possible values associated with those arcs

RDF Node

```
:alice schema:name   "Alice";
       schema:knows  :bob .
```

Abstract shape of a node that represents a User

```
IRI schema:name   string  1
    schema:knows  IRI     0, 1,...
```

ShEx

```
<UserShape> IRI {
  schema:name   xsd:string    ;
  schema:knows  IRI           *
}
```

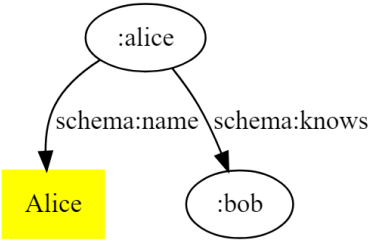

Understanding the problem

RDF flexibility

- Mixed use of objects & literals
- Example:
  - Values of `schema:creator` can be:
    - **`string`** or `schema:Person`
  - in the same data

```
:angie schema:creator "Keith Richards" ;
       schema:creator [
         schema:firstName "Mick" ;
         schema:lastName  "Jagger"
       ] .
```

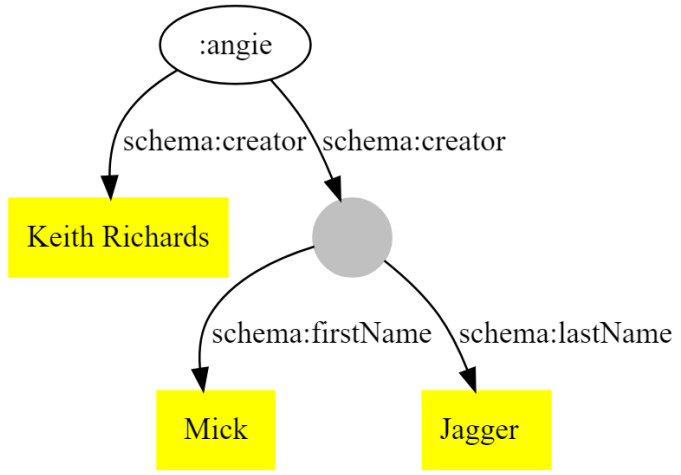

Lots of examples at http://schema.org

Understanding the problem

Repeated properties

The same property can be used for different purposes in the same data

Example: A product must have 2 codes with different structure

```
:product schema:productID "isbn:123-456-789";
         schema:productID "code456" .
```

A practical example from FHIR

See: http://hl7-fhir.github.io/observation-example-bloodpressure.ttl.html

Understanding the problem

- Shapes ≠ types
- Nodes in RDF graphs can have zero, one or many `rdf:type` declarations
  - One type can be used for multiple purposes (`schema:Person`)
  - Data doesn't need to be annotated with fully discriminating types
    - Nodes with type `schema:Person` can also be customers, patients, etc...
    - Different meanings and different structure depending on the context
    - Different validation constraints at different contexts

Shapes vs Ontology

Ontologies ≠ Shapes ≠ instance data

- Ontologies are usually focused on domain entities (higher level)
- RDF validation/shapes focused on RDF graph features (lower level)

Different levels

Ontology

```
:Person a owl:Class ;
   rdfs:subClassOf [a owl:Restriction ;
                    owl:onProperty  :hasParent ;
                    owl:qualifiedCardinality 2 ;
                    owl:onClass :Person ].
```

Shapes
RDF *Validation*
Constraints

```
<PersonShape> IRI {
 :hasParent @<PersonShape> {0,2}
}
```

Instance data

```
:alice :hasParent :bob, :carol .
:bob   :hasParent  :dave .
```

Previous RDF validation approaches

- SPARQL based
  - Plain SPARQL
  - SPIN: http://spinrdf.org/
- OWL based
  - Stardog ICV
    - http://docs.stardog.com/icv/icv-specification.html
- Grammar based
  - OSLC Resource Shapes
    - https://www.w3.org/Submission/2014/SUBM-shapes-20140211/

Define SPARQL queries that detect errors

Pros:

- Expressive
- Ubiquitous

Cons

- Expressive
- Idiomatic - many ways to encode the same constraint

Example: SPARQL query to check that…

There is one `schema:name` which must be a `xsd:string` and one `schema:gender` must be `schema:Male` or `schema:Female`

```
ASK {{ SELECT ?Person {
      ?Person schema:name ?o .
    } GROUP BY ?Person HAVING (COUNT(*)=1)
  }
  { SELECT ?Person {
      ?Person schema:name ?o .
      FILTER ( isLiteral(?o) &&
               datatype(?o) = xsd:string )
     } GROUP BY ?Person HAVING (COUNT(*)=1)
  }
  { SELECT ?Person (COUNT(*) AS ?c1) {
      ?Person schema:gender ?o .
    } GROUP BY ?Person HAVING (COUNT(*)=1)}
    { SELECT ?Person (COUNT(*) AS ?c2) {
      ?S schema:gender ?o .
      FILTER ((?o = schema:Female ||
               ?o = schema:Male))
    } GROUP BY ?Person HAVING (COUNT(*)=1)}
    FILTER (?c1 = ?c2)
}
```

SPIN

SPARQL inferencing notation http://spinrdf.org/

- Developed by TopQuadrant
- Commercial product

Vocabulary associated with user-defined functions in SPARQL

SPIN has influenced SHACL (see later)

Stardog ICV

ICV - Integrity Constraint Validation

- Commercial product

OWL with unique name assumption and closed world

Compiled to SPARQL

More info: http://docs.stardog.com/icv/icv-specification.html

OSLC Resource Shapes

OSLC Resource Shapes

https://www.w3.org/Submission/shapes/

- Grammar based approach
- Language for RDF validation
- Input for ShEx and SHACL

```
:user a rs:ResourceShape ;
 rs:property [
  rs:name "name" ;
  rs:propertyDefinition schema:name ;
  rs:valueType xsd:string ;
  rs:occurs rs:Exactly-one ;
] ;
 rs:property [
  rs:name "gender" ;
  rs:propertyDefinition schema:gender ;
  rs:allowedValue schema:Male, schema:Female ;
  rs:occurs rs:Zero-or-one ;
].
```

Other approaches

Dublin Core Application profiles (K. Coyle, T. Baker)

http://dublincore.org/documents/dc-dsp/

RDF Data Descriptions (Fischer et al)

http://ceur-ws.org/Vol-1330/paper-33.pdf

RDFUnit (D. Kontokostas)

http://aksw.org/Projects/RDFUnit.html

...

ShEx and SHACL

2013 RDF Validation Workshop

- Conclusions of the workshop:
  - *There is a need of a higher level, concise language for RDF Validation*
- ShEx initially proposed (v 1.0)

2014 W3c Data Shapes WG chartered

2017 SHACL accepted as W3C recommendation

2017 ShEx 2.0 released as W3C Community group draft

2019 ShEx adopted by Wikidata

Continue this tutorial with...

ShEx by example 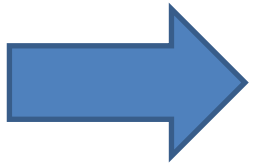 https://doi.org/10.6084/m9.figshare.13174571

SHACL by example 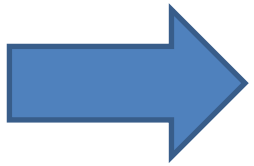 https://doi.org/10.6084/m9.figshare.13174577

ShEx and SHACL compared 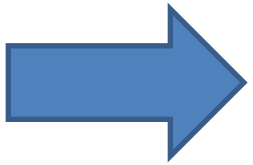 https://doi.org/10.6084/m9.figshare.13174583

Shapes applications and tools 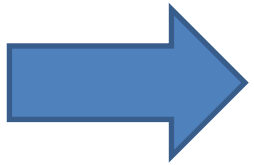 https://doi.org/10.6084/m9.figshare.13174586